



A Practical guide to HPC

mattia.pellegrino@unipr.it

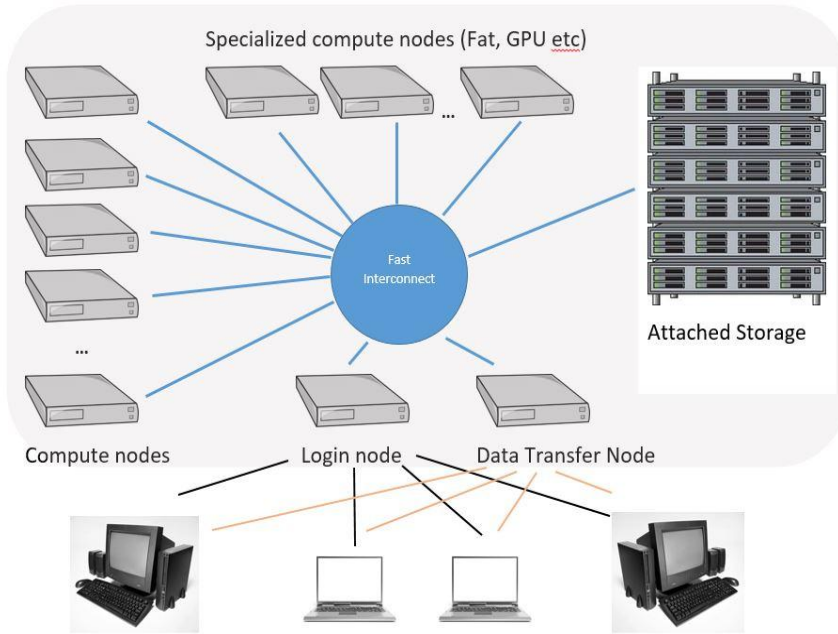
What is a cluster? (1)

- A **computer cluster** is a group of tightly coupled computers (**nodes**) that work together closely
- They can be seen as a **single computer**
- Clusters are commonly connected through **fast local area network**
- Clusters have evolved to support various applications:
 - **Ecommerce**
 - **High performance database** applications
- Clusters are usually deployed to improve **speed** and **reliability** over that provided by a single computer

What is a cluster? (2)

- In cluster computing each node (within a cluster) is an **independent system**
 - Has its own **operating system**
 - **Private** memory
 - Its own **file system** (in some cases)
- Programs or software run on clusters usually employ a procedure called “**message passing**” (memory is not shared)
- Cluster computing can also be used as a **low-cost** form of **parallel processing** for scientific applications

What is a cluster? (3)



History

- The first clustering product was **ARCnet**, distributed and developed by Datapoint in **1977**
- Then in **1980 VAXcluster** produced a new product
- **Sun Microsystem, Microsoft** and other leading hardware and software companies offer clustering packages

Serial computing

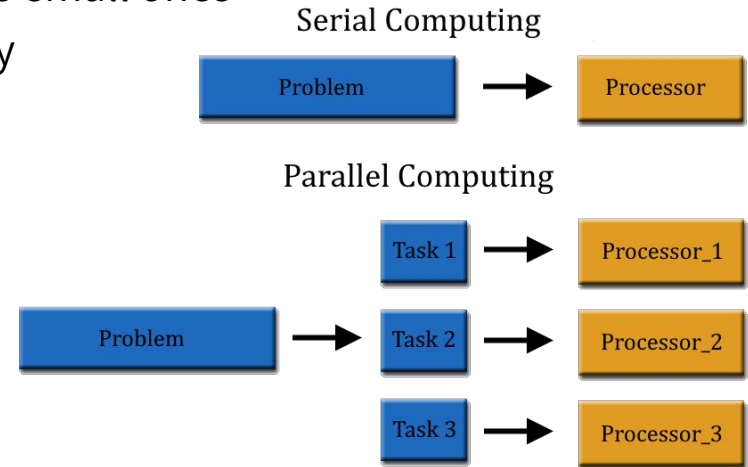
- **Serial computing**

- Based on **Von Neumann architecture**, it have to process data and instructions respecting a sequential order
- Limits:
 - Need **high speed** cpu
 - **Light** speed
 - **Heat**
 - **VLSI** (very large scale integration) is not infinite

Parallel computing


- **Parallel computing**

- “*Big problems can be often divided into small ones*”
- Problems can be solved simultaneously
- Advantages:
 - Improve computing speed
 - Break memory limits
 - Exploit non-local resources



Distributed Memory

- Distributed memory requires a communication network for the information exchange:
 - Each **processor** has its own **local memory**.
 - Each **memory** has a **separate, independent** address space.
 - **Read/write** operations are local
 - To allow a task to **access remote data**, the programmer must explicitly **manage** the **communication** among tasks.
 - The corresponding programming model is called **message passing**




MPI

- **Message passing Interface:**
 - Communication protocol for **message exchanges**
 - There are **multiple** implementations of MPI
 - It is **not** a **language**
 - The standard has been defined through an **open process** by **community** of parallel computing vendor, computer scientists, and application developers
 - It does not have a debugging facility and this makes MPI **lightweight**
 - Several MPI implementations: **MPICH**, **Open MPI** (best overall), **Intel MPI Library**

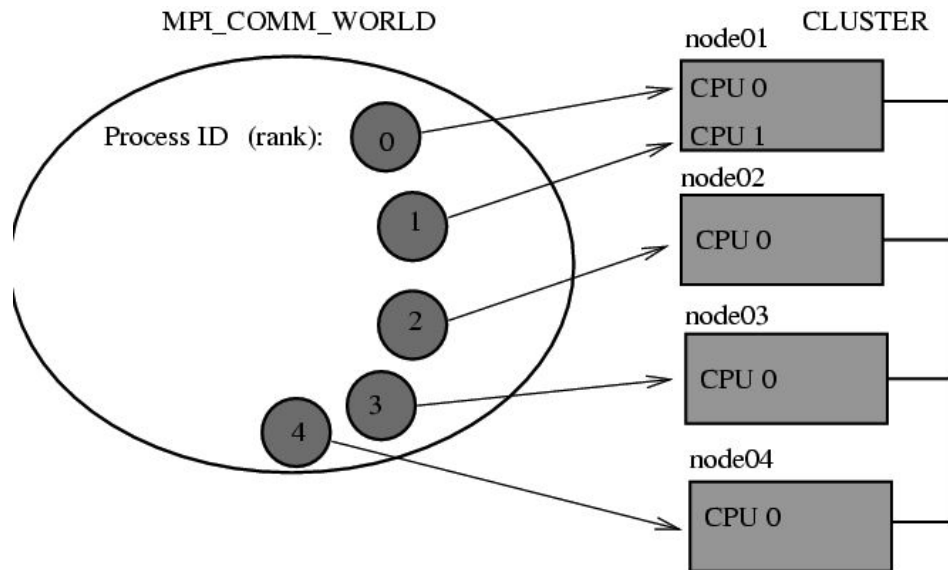
MPI

- Cooperation among processes is based on **explicit communication**
- A **sender** process and a **receiver** process exchange messages
- Each process is an **instance** of running sub-program.
- Each process is identified by an integer number, called **rank**, ranging from **0** to **n-1**, where n (size) is the total number of processes.



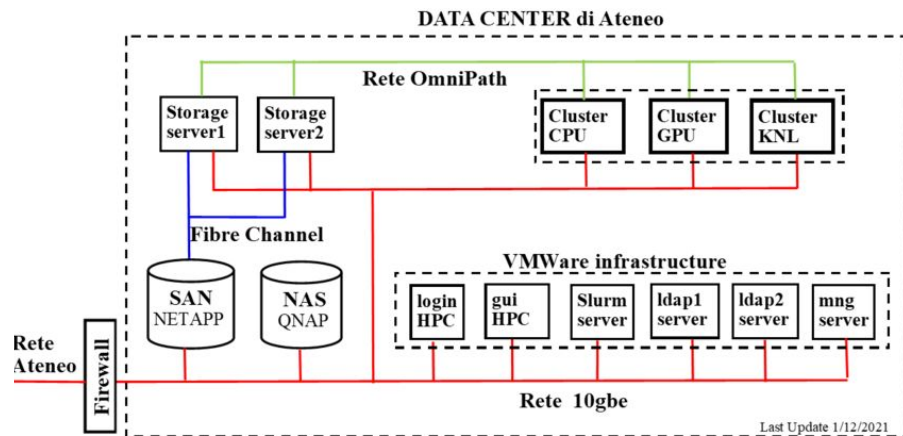
MPI - SPECIFICATIONS

- MPI allows you to create logical **groups** of processes
 - In each **group**, a process is identified by its **rank**.
- **Communicators** are **objects** that handle communication between processes.
 - **intra-communicator** handles processes within a single group
 - **inter-communicator** handles communication between two distinct groups.
- By default, there is a single group that contains all your processes, and the intra-communicator **MPI_COMM_WORLD**



Unipr - HPC

- **Website:**
 - <https://www.hpc.unipr.it>
- **Project Description:**
 - <https://www.hpc.unipr.it/dokuwiki/doku.php?id=calcoloscientifico:progetto>
- **User Guide:**
 - <https://www.hpc.unipr.it/dokuwiki/doku.php?id=calcoloscientifico:userguide>



Unipr - HPC - Nodes

Partition	Node Name	CPU Type	HT	#Cores	MEM (GB)	GPU	Total	Owner
bdw, cpu	wn01-wn08	2 INTEL XEON E5-2683v4 2.1GHz 16c	NO	32	128	0	256 cores	Public
bdw, cpu	wn09-wn17	2 INTEL XEON E5-2680v4 2.4GHz 14c	NO	28	128	0	252 cores	Public
bdw, cpu	wn33	2 INTEL XEON E5-2683v4 2.1GHz 16c	NO	32	1024	0	32 cores	Public
bdw, cpu	wn34	4 INTEL XEON E7-8880v4 2.2GHz 22c	NO	88	512	0	88 cores	Public
cpu	wn35:wn36	4 INTEL XEON E5-6252n 2.3 GHz 24c	NO	96	512	0	192 cores	Public
cpu_infn	wn22	4 INTEL XEON E5-5218 2.3 GHz 16c	NO	64	384	0	64 cores	Private
cpu_bioscienze	wn23	4 INTEL XEON E5-5218 2.3 GHz 16c	NO	64	384	0	64 cores	Private
cpu_mmm	wn24-wn25	4 INTEL XEON E5-5218 2.3 GHz 16c	NO	64	384	0	128 cores	Private
cpu_guest	wn22-wn25	4 INTEL XEON E5-5218 2.3 GHz 16c	NO	64	384	0	256 cores	Guest
gpu	wn41-wn42	2 INTEL XEON E5-2683v4 2.1GHz 16c	NO	32	128	7 P100	14 GPU	Public
gpu	wn44	2 AMD EPYC 7352 2.3GHz 24c	NO	32	512	4 A100	4 GPU	Public
gpu_hylab	wn43	2 INTEL XEON Silver 4210 2.2GHz 10c	NO	20	384	2 V100	2 GPU	Private
gpu_guest	wn43	2 INTEL XEON Silver 4210 2.2GHz 10c	NO	20	384	2 V100	2 GPU	Guest
kn1	wn51-wn54	1 INTEL XEON PHI 7250 1.4GHz 68c	YES	272	192	0	1088 cores	Public
skl, cpu	wn18-wn19	4 INTEL XEON E5-6140 2.3GHz 18c	NO	72	384	0	144 cores	Public
skl_mm1	wn20	4 INTEL XEON E5-6140 2.3GHz 18c	NO	72	384	0	72 cores	Private
skl_infn, cpu_infn	wn21	4 INTEL XEON E5-6140 2.3GHz 18c	NO	72	384	0	72 cores	Private
skl_guest, cpu_guest	wn20-wn21	4 INTEL XEON E5-6140 2.3GHz 18c	NO	72	384	0	144 cores	Guest
vrt	wn61-wn64	2 INTEL XEON E5-2620v4 2.1GHz 8c	YES	8	64	0	32 cores	Public

GPUs

Node Name	GPU Type
wn41-wn42	NVIDIA Corporation GP100GL [Tesla P100 PCIe 12GB] (rev a1)
wn43	NVIDIA Corporation GV100GL [Tesla V100 PCIe 32GB] (rev a1)
wn44	NVIDIA Corporation GA100GL [Tesla A100 PCIe 40GB] (rev a1)

Unipr - HPC - Limits

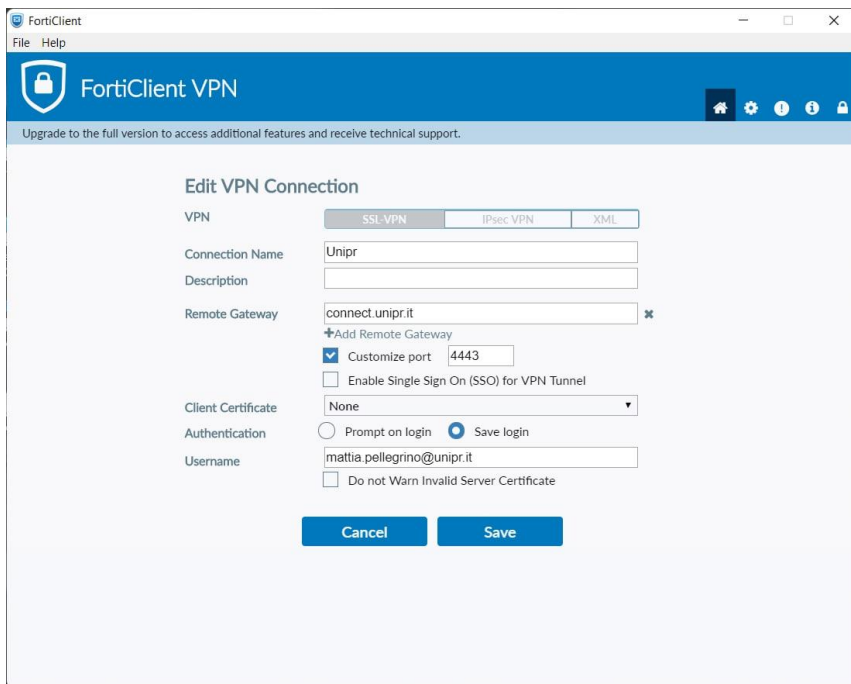
Partition	Nodes	Job Resources	TimeLimit	Max Running	Max Submit	Max nodes	Audience
				per user	per user	per job	
cpu	wn01:19,wn33:36		3-00:00:00	24	2000	8	Public
cpu_mmm	wn24:25		3-00:00:00				Private
cpu_bioscienze	wn23		3-00:00:00				Private
cpu_infn	wn21:22		3-00:00:00				Private
cpu_guest	wn20:25		3-00:00:00				Guest
bdw	wn01:08,wn09:17	2-628 cores	3-00:00:00	24	2000	8	Public
skl	wn18,wn19	2-124 cores	3-00:00:00	24	2000	8	Public
skl_mm1	wn20	2-72 cores	3-00:00:00				Private
skl_infn	wn21		3-00:00:00				Private
skl_guest	wn20:21		1-00:00:00				Guest
kn1	wn51:54		5-00:00:00	16	2000		Public
gpu	wn41:42,wn44		0-24:00:00	6	2000		Public
gpu_hylab	wn43	1-2 gpu	0-24:00:00				Private
gpu_guest	wn43	1-2 gpu	0-24:00:00				Guest
vrt	wn61:64	1core	10-00:00:00	24	2000	1	Public
mngr	Reserved						Manage

SLURM

- Slurm is an open source, fault-tolerant, and highly scalable **cluster management** and **job scheduling system** for large and small Linux clusters
- SLURM manages user jobs with the following key characteristics:
 - set of requested resources:
 - X number of computing resources: **nodes** (including all their CPUs and cores) or **CPUs** (including all their cores) or cores
 - X number of accelerators (**GPUs**)
 - X **amount** of **memory**: either per node or per (logical) CPU
 - X the (wall)**time** needed for the user's tasks to complete their work
 - a set of **constraints** limiting jobs to nodes with specific features
 - a requested **node partition** (job queue)

PRACTICAL GUIDE

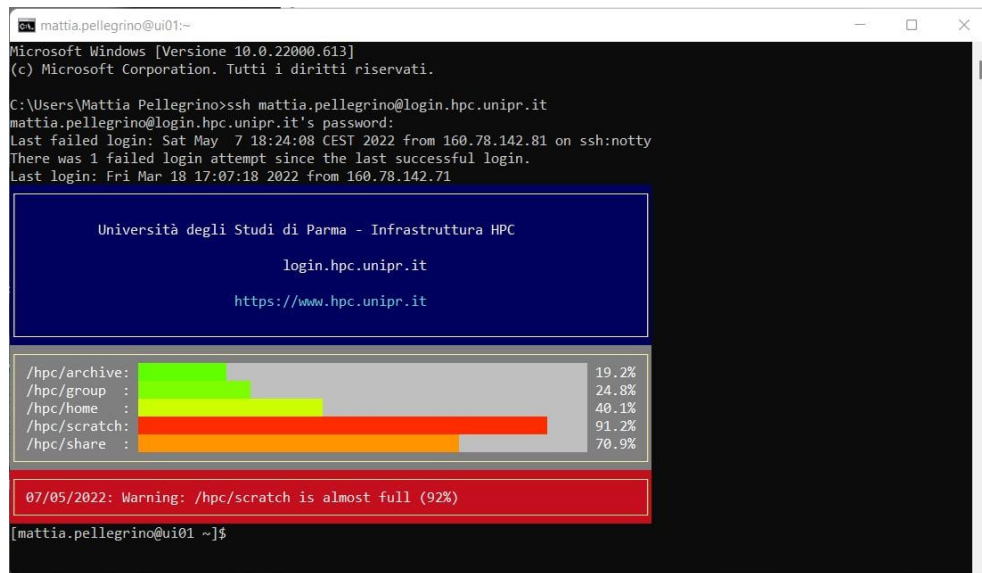
- First of all, you need to be connected to UNIPR wireless network, instead you must use a VPN in order to connect to the HPC facility



- Steps:**
 - Download **FortiClient**:
<https://www.fortinet.com/it>
 - Use this parameters to setting a new VPN connection:
 - Host: **connect.unipr.it**
 - Port: **4443**
 - Use you UNIPR credentials to login

PRACTICAL GUIDE

- In order to connect to the HPC Cluster, you must use an SSH protocol connection



```
mattia.pellegrino@ui01:~  
Microsoft Windows [Versione 10.0.22000.613]  
(c) Microsoft Corporation. Tutti i diritti riservati.  
  
C:\Users\Mattia Pellegrino>ssh mattia.pellegrino@login.hpc.unipr.it  
mattia.pellegrino@login.hpc.unipr.it's password:  
Last failed login: Sat May 7 18:24:08 CEST 2022 from 160.78.142.81 on ssh:notty  
There was 1 failed login attempt since the last successful login.  
Last login: Fri Mar 18 17:07:18 2022 from 160.78.142.71  
  
Università degli Studi di Parma - Infrastruttura HPC  
  
login.hpc.unipr.it  
  
https://www.hpc.unipr.it  
  
/hpc/archive: 19.2%  
/hpc/group : 24.8%  
/hpc/home : 40.1%  
/hpc/scratch: 91.2%  
/hpc/share : 70.9%  
  
07/05/2022: Warning: /hpc/scratch is almost full (92%)  
[mattia.pellegrino@ui01 ~]$
```

- **Steps:**
 - Open the system terminal
 - **ssh <nomeutente>@login.hpc.unipr.it**
 - Password: your institutional password

FILE TRANSFER

- **SSH** is the only protocol for external communication and can also be used for file transfer.
- If you use a Unix-like client (Linux, MacOS X) you can use the command **scp** or **sftp**.
- On Windows systems, the most used tool is WinSCP (<https://winscp.net/eng/docs/introduction>).

scp example: **scp /path/to/local/file remote_user@remote_host:/path/to/remote/file**

scp foo.file mattia.pellegrino@login.hpc.unipr.it:/hpc/group/G_SOWIDE/10-05-2022/python/NN/foo.file

PRACTICAL GUIDE

- To correctly launch a job on hpc you have to write an .sh script. This script is required to use SLURM functionalities

This is not a comment!

```
1 #!/bin/bash
2
3 #SBATCH --partition=bdw
4 #SBATCH --output=%x.o%j
5 #SBATCH --nodes=1
6 #SBATCH --ntasks=32
7 #SBATCH --cpus-per-task=1
8
9 #SBATCH --mem=900G
10 ##SBATCH --mem-per-cpu=27G
11
12 #SBATCH --time=3-00:00:00
13
14 module load java/jdk/12.0.1
15
16 JOBNAME=$SLURM_JOB_NAME # use the job-name specified above
17
18 # Run 1 job per task
19 N_JOB=$SLURM_NTASKS # create as many jobs as tasks
20
21 for((i=1;i<=$N_JOB;i++))
22 do
23     if [ $i -eq 1 ] || [ $i -eq $N_JOB ]
24     then
25         if [ $i -eq 1 ]
26         then
27             java -jar powerlaw_manualip.jar b tcp://$HOSTNAME -ib:61616 &
28             sleep 20
29         else
30             java -jar powerlaw_manualip.jar tcp://$HOSTNAME -ib:61616 i
31         fi
32     else
33         java -jar powerlaw_manualip.jar tcp://$HOSTNAME -ib:61616 n &
34         sleep 5
35     fi
36 done
37
38 #Wait for all
39 wait
40
41 echo
42 echo "All done. Checking results:"
43 grep "PI" $JOBNAME.*/log
44
45
```

```
1 #!/bin/bash
2
3 #SBATCH --partition=bdw
4 #SBATCH --output=%x.o%j
5 #SBATCH --nodes=1
6 #SBATCH --ntasks=32
7 #SBATCH --cpus-per-task=1
8
9 #SBATCH --mem=900G
10 ##SBATCH --mem-per-cpu=27G
11
12 #SBATCH --time=3-00:00:00
13
14 module load java/jdk/12.0.1
15
```

#partition name
Standard output and error log
Run all processes on a single node
Run a single task
Number of CPU cores per task

Total memory limit
→ comment

Time limit hrs:min:sec

#load the desire module
module avail to see all module supported by the HPC

How to launch a script: sbatch <nomescript>.sh

Queue: queue | grep <partition_name>

queue | grep <username>

queue | grep <job_id>

Delete a job: scancel <job_id>


Job's details: scontrol show jobid <job_id>

ACTODES

- Java **framework**
- Support the **actor** paradigm:
 - Simple and independent **agents**
 - Based on **message queue** and **exchange**
- It can be used in **distributed architecture**

To ignore time limit!

JAVA ON HPC (actodes)



activemq.xml



mpiLaunch.sh



MultiActorS...



test.sh

folder

```
1 #!/bin/bash
2 #SBATCH --job-name=test
3 #SBATCH --output=%x.o%j
4 #SBATCH --error=%x.e%j
5 #SBATCH --nodes=2
6 #SBATCH --ntasks-per-node=5
7 #SBATCH --partition=bdw
8 #SBATCH --mem=200G
9 ##SBATCH --odelist=wn33
10 ##SBATCH --mem-per-cpu=15G
11 #SBATCH --cpus-per-task=1
12 #SBATCH --time=0-02:00:00
13 #SBATCH --account=g_sowide
14
15 module load gnu openmpi
16 module load java/jdk/12.0.1
17
18 mpirun --mca btl_tcp_if_include ib0 ./test.sh
```


InfiniBand (computer networking communication standard)

Modular component architecture

Use TCP protocol

```
1 #!/usr/bin/env bash
2
3 test -n "$SLURM_JOB_NAME" || exit
4
5 log="${SLURM_JOB_NAME}.${SLURM_JOB_ID}.${OMPI_COMM_WORLD_RANK}.${hostname -s}.log"
6
7 echo "This is rank $OMPI_COMM_WORLD_RANK on $(hostname -s)." >> "$log"
8 echo "The master node is ${HOSTNAME}." >> "$log"
9 echo "Host $(hostname -ib)." >> "$log"
10
11 sleep $((2*(OMPI_COMM_WORLD_RANK+1)))
12
13 echo "OK from rank $OMPI_COMM_WORLD_RANK"
14
15 # Run 1 job per task
16 N_JOB=$OMPI_COMM_WORLD_RANK # create as many jobs as tasks
17 N_TASK=$OMPI_COMM_WORLD_LOCAL_SIZE
18 N_JOB_MAX=$((OMPI_COMM_WORLD_SIZE - 1))
19
20 echo $N_JOB
21 echo $N_TASK
22 echo $N_JOB_MAX
23
24 if [ $N_JOB -gt 0 ]
25 then
26     echo "Need to Sleep $((5 * $N_JOB)) sec"
27     sleep $((5 * $N_JOB))
28 fi
29
30 if [ $N_JOB -eq 0 ] || [ $N_JOB -eq $N_JOB_MAX ]
31 then
32     if [ $N_JOB -eq 0 ]
33     then
34         echo "Lancio del Broker"
35         java -jar MultiActorSpaces.jar b tcp://$HOSTNAME-ib:61616 >> "$log"
36     else
37         echo "Lancio dell'Initiator"
38         java -jar MultiActorSpaces.jar i tcp://$HOSTNAME-ib:61616 >> "$log"
39     fi
40 else
41     java -jar MultiActorSpaces.jar n tcp://$HOSTNAME-ib:61616 >> "$log"
42 fi
```

ACTODES ON HPC



PYTHON for ML



cars.csv



cars_regres...



train.sh

You can select a specific GPU

Load ML resources


Release ML resources

```
train.sh
1  #!/bin/bash
2
3  #SBATCH --partition=gpu
4  #SBATCH --nodes=1
5  #SBATCH --ntasks-per-node=1
6  #SBATCH --gres=gpu:a100:1
7  #SBATCH --time 0-01:00:00
8  #SBATCH --mem=1G
9
10 echo $SLURM_JOB_NODELIST
11
12 echo #OMP_NUM_THREADS : $OMP_NUM_THREADS
13
14 module load miniconda3
15 source "$CONDA_PREFIX/etc/profile.d/conda.sh"
16 conda activate machine-learning-cuda-10.2
17 #pip install einops
18
19
20 python cars_regression.py
21
22 conda deactivate
```

GOSSIP ALGORITHM

Exercise

- Systems-oriented **computational and communication paradigm** distributed on a **large scale** with high dynamic characteristics
- Probabilistic approach
- Main features:
 - **Simplicity**
 - **Scalability**
 - **Efficiency**
 - **Robustness**



EXERCISE

A Gossip Algorithm

- Modify the source code available at the following address:
https://github.com/sowide/HPC_practice
- Each agent involved in the simulation can be: **infected or healthy**
- Agents meet each other **randomly**
- Every time an **interaction** occurs, there is a chance that an **infected** agent **will infect** a **healthy** one

EXERCISE

- Classes and function to complete:
 - **Phases.class**
 - **PersonManager.class:**
 - BuildPopulation(), infectionInPartition(), setinfectedpeople(), saveMeasure()
 - **Manager.class**
 - messageDataHandler()
- When the simulator is ready (**test it locally with 100K agents**) launch the simulator on the HPC **with 1M of agents**, utilizing **8,16,24,32** cores. Finally, measure the temporal performance
- **Extra:** Add 1 or more personalized phases (you're free to decide which phases you want)
- Deliver source code: <https://shorturl.at/bfBJR> [name_surname.zip]

EXERCISE

Login

- **Login and take reservation:**

- `ssh <nomeutente>@login.hpc.unipr.it`
- `ssh-keygen -t rsa -P "" -f ~/.ssh/id_rsa`
- `cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys`
- `chmod 0600 ~/.ssh/authorized_keys`
- `newgrp T_2022_SISTEMI_DISTRIBUITI`
- **Aggiungere allo script mpiLunch.sh:**
 - `#SBATCH --reservation=t_2022_sistemi_distribuiti-20230418`
 - `#SBATCH --account=t_2022_sistemi_distribuiti`

Local Execution

Initiator Class

```
switch (s)
{
    //lancio del broker
    case "b":
        c.setExecutor(new CycleScheduler(TimeoutMeasure.CY));
        c.setConnector(new ActiveMqConnector(s2, "")); // --> c.setConnector(new ActiveMqConnector(true));
        c.addService(new Creator());
        break;
    //lancio di un generico nodo
    case "n":
        c.setExecutor(new CycleScheduler(TimeoutMeasure.CY));
        c.setConnector(new ActiveMqConnector(s2)); // --> c.setConnector(new ActiveMqConnector(false));
        c.addService(new Creator());
        break;
    //lancio dell'initiator
    case "i":
        c.setExecutor(new CycleScheduler(
            new PersonManager(true), TimeoutMeasure.CY));
        c.setConnector(new ActiveMqConnector(s2)); // --> c.setConnector(new ActiveMqConnector(false));
        break;
    default:
        c.setExecutor(new CycleScheduler(
            new PersonManager(true), TimeoutMeasure.CY));
}

c.start();
```

s2 is the string that contains the IP address of the broker:

tcp://x.x.x.x:port

(es. tcp://127.0.0.1:61616)

- To run the simulator on local machine you can also change the object argument **ActiveMqConnector**
- To launch the simulator correctly you need to:
 - Launch a broker instance **(b)**
 - Launch one or more generic instances **(n)**
 - Launch an initiator instance to start the simulation (master) **(i)**